

CS168 Spring Assignment 4
SUNet ID(s): 05794739
Name(s): Luis A. Perez
Collaborators:

By turning in this assignment, I agree by the Stanford honor code and declare that all of this is my own work.

Part 1

- (a) See appendix for code.
- (b) PCA would recover a uniformly random slope. There is no axis which maximizes the variance of the data, as such the recovered best-fit line would be random.

LS will recover a line with near 0.0 slope (and y-intercept of 0.5), since this will minimize the distance to the points in the $[0,1] \times [0,1]$ square. Any other slope will tend to increase the distance to the opposing corners of the square.

See code in appendix for simulation.

- (c) The generated plot is seen in Figure 1




Figure 1: PCA and Least Squares best-fit line slopes for varying levels of noise in measurement but not input signal. True slope of data is 2.0

- (d) The generated plot is seen in Figure 2
- (e) PCA does poorly when the noise exists only in Y because it tries to find the direction in which the variability of the data is maximized. By adding noise (additional variability)

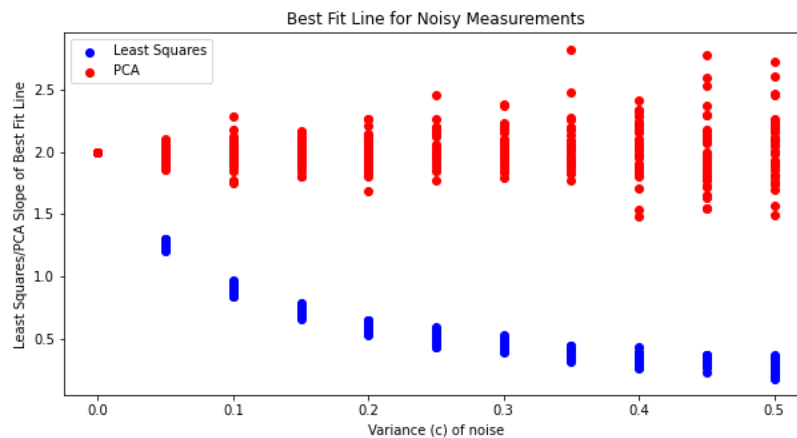


Figure 2: PCA and Least Squares best-fit line slopes for varying levels of noise in measurement and input signal. True slope of data is 2.0

along the Y -axis only, PCA accounts for this by shifting the slope of the line closer to vertical.

On the other hand, when there's noise along both X and Y , the variance along both axis sort of cancels itself out. As such, the principal component will be in the direction of the original slope.

In this situation where both X and Y are noisy, however, LS will do poorly. Drawing an analogy to our uniformly random points in the unit square, by adding noise along both axes we're approximating a more uniform distribution. As such, the LS solution will begin to favor lines with 0 slope, since this minimizes the distance to all points.

Part 2

- (a) The dimension would be 10,101, since this is the number of features.
- (b) The scatter plot is shown in Figure 3.

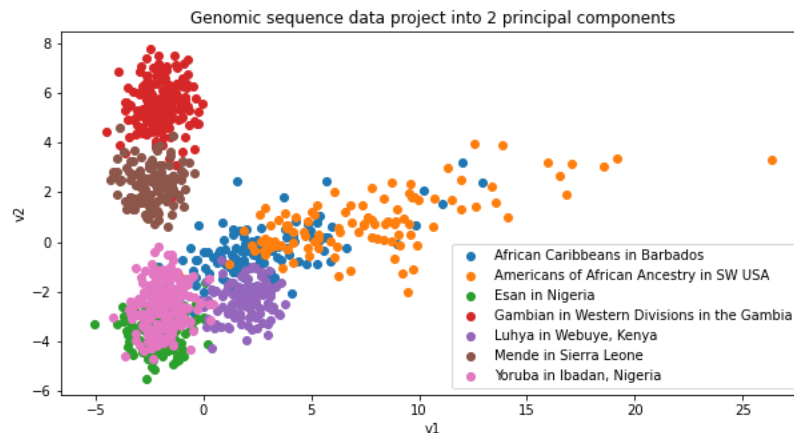


Figure 3: Genomic data for 995 individuals projected onto the first two principal components as per PCA method.

- (c) From the plot in Figure 3, we note the following.
 Most populations tends to cluster in groups, with Gambian, Mende, Yoruba, and Esan populations varying primarily along the y-axis, while Luhya, African Caribbeans, and African American's varying the most along the x-axis.
 As such, it appears that v_1 (the principal component) has a temporal aspect, since African Caribbeans and African Americans are more recent populations and these vary the most along this axis. It seems to measure the amount of migration. v_2 on the other hand appears to capture primarily differences within African-continent populations, so it likely captures more historical/geographical distinctions between these native populations.
- (d) The scatter plot is shown in Figure 4.
- (e) The third principal component appears to capture gender.
- (f) We see the specified plot in Figure 5. From this plot, we can see that the last few hundred nucleobases appears to encode information about the Y chromosome, whereas the begining few nucleobases appears to encode information about the X chromosome (larger).

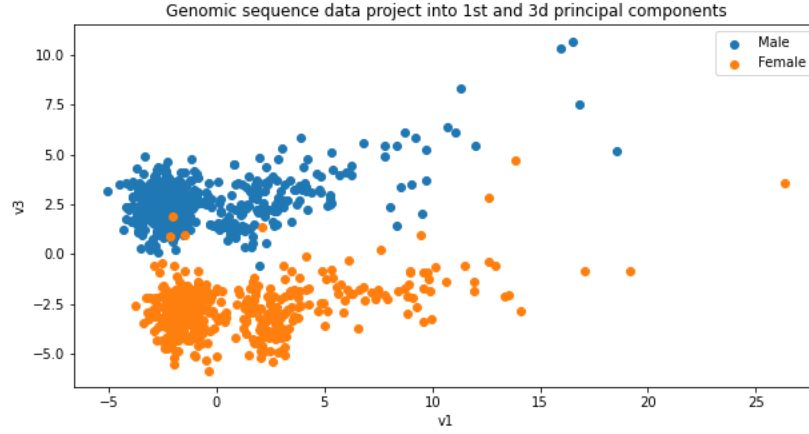


Figure 4: Genomic data for 995 individuals projected onto the first and third principal components as per PCA method.

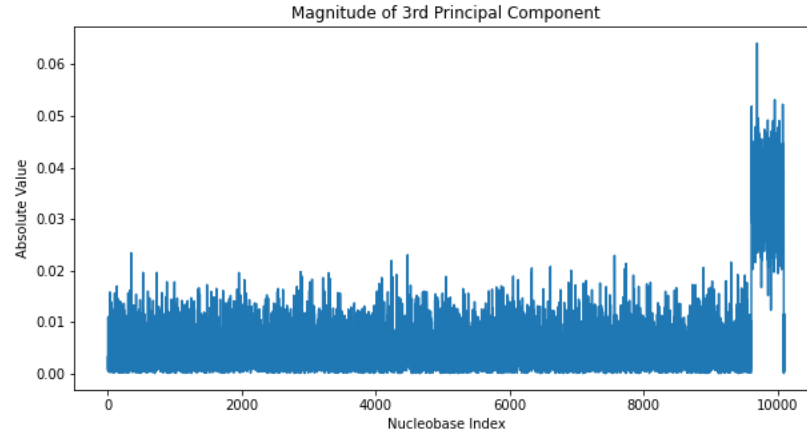


Figure 5: Nucleobase plot of the third principal component which we hypothesized encodes gender.

- (g) We create a bijection by mapping each of A, T, C, G to a one-hot encoded vector. What we mean by this is that our original matrix $X \in \mathbb{R}^{n \times d}$ where d is the feature dimension, would instead become $X' \in \mathbb{R}^{n \times (dx4)}$ where we conceptually group each of the 4 dimensions to a binary label.
- (h) We recreate the plot from (b) above in Figure 6. We do not see any immediate added value.
- (i) It appears that the fourth component capture familiar relationship. We see two cluster along the y-axis in Figure 7, where all unrelated individuals are separated into their own own, while individuals from the same family are all clustered together.

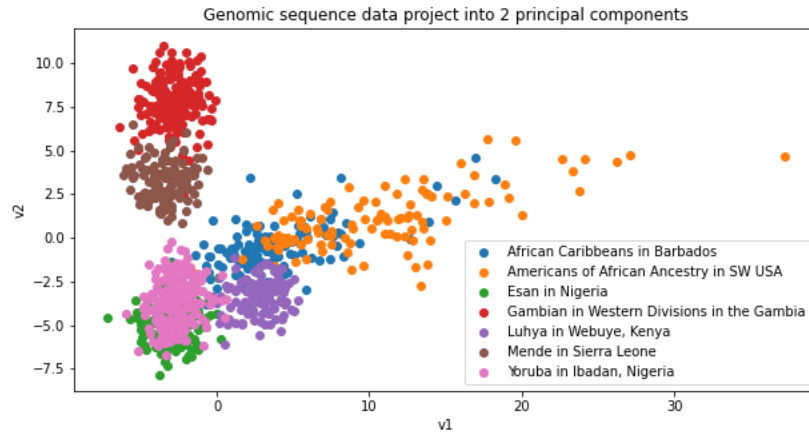


Figure 6: Genomic data for 995 individuals projected onto the first two principal components as per PCA method with one-hot encoded nucleobases.

As such, this seems to encode some sense of relatedness to each other in the dataset.

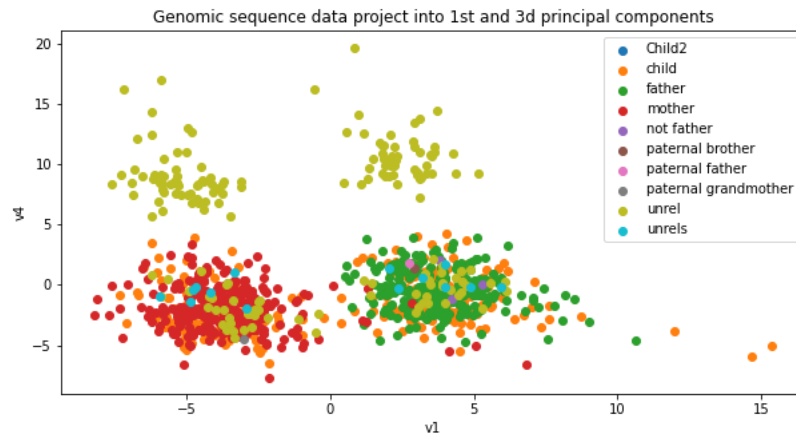


Figure 7: Genomic data for 995 individuals projected onto the third (x-axis) and fourth (y-axis) principal components.

(j) NOT DONE.

Appendix

HW4

April 30, 2020

1 CS 168 Spring Assignment 4

SUNet ID(s): 05794739

Name(s): Luis A. Perez

Collaborators: None

By turning in this assignment, I agree by the Stanford honor code and declare that all of this is my own work.

2 Imports

```
[180]: import collections
import matplotlib.pyplot as plt
import scipy

import numpy as np
from scipy import linalg
from sklearn import decomposition
import pandas as pd
import seaborn as sns
import os
import warnings

from typing import Dict, List, Text, Tuple

# Make figure larger
plt.rcParams['figure.figsize'] = [10, 5]

# Set numpy seed for consistent results.
np.random.seed(1)
```

3 Part 1

3.1 Part 1a

```
[34]: def pca_recover(X, Y):  
    """Returns the slope of the first component of PCA applied to [X, Y].  
  
    Args:  
        X: An (n,1) matrix of points.  
        Y: An (n,1) matrix of points.  
    """  
    data = np.concatenate((X,Y), axis=1)  
    data = data - np.mean(data, axis=0)  
    A = np.dot(data.T, data)  
    # Compute largest eigenvector.  
    _, v = linalg.eigh(A, eigvals=(1,1))  
    v = v.flatten()  
    return v[1] / v[0]
```

```
[38]: def ls_recover(X, Y):  
    """Returns the slop of the best fit line as per least squares. See ↵  
    ↪pca_recover()."""  
    centeredX, centeredY = X - np.mean(X, axis=0), Y - np.mean(Y, axis=0)  
    return (np.dot(centeredX.T, centeredY) / np.sum(centeredX**2)).flatten()[0]
```

```
[39]: def problem1a():  
    X = np.arange(0.001, 1.001, 0.001)  
    X = np.reshape(X, (X.shape[0], 1))  
    Y = 2*X  
    return pca_recover(X,Y), ls_recover(X, Y)
```

```
[40]: problem1a()
```

```
[40]: (2.0, 2.0)
```

```
[105]: def problem1b():  
    pca, ls = [], []  
    for _ in range(100000):  
        X = np.random.uniform(size=(1000, 1))  
        Y = np.random.uniform(size=(1000, 1))  
        pca.append(pca_recover(X,Y))  
        ls.append(ls_recover(X, Y))  
    return np.mean(pca), np.mean(ls), np.std(pca), np.std(ls)
```

```
[106]: problem1b()
```



```
[106]: (-0.2550136867924121,  
        3.252073550279325e-05,  
        142.44712135685342,  
        0.03173951330042009)
```

```
[277]: def problem1cd(title, add_x_noise):  
        X = np.arange(0.001, 1.001, 0.001)  
        X = np.reshape(X, (X.shape[0], 1))  
        cs = np.arange(0, .55, 0.05)  
        ls = []  
        pca = []  
        for _ in range(30):  
            for c in cs:  
                if add_x_noise:  
                    noisyX = X + np.random.normal(scale=np.sqrt(c), size=X.shape)  
                else:  
                    noisyX = X  
                Y = 2*X + np.random.normal(scale=np.sqrt(c), size=X.shape)  
                pca.append((c, pca_recover(noisyX, Y)))  
                ls.append((c, ls_recover(noisyX, Y)))  
  
        x, y = zip(*ls)  
        plt.scatter(x, y, label="Least Squares", color="b")  
        x, y = zip(*pca)  
        plt.scatter(x, y, label="PCA", color="r")  
        plt.title("Best Fit Line for Noisy Measurements")  
        plt.xlabel("Variance (c) of noise")  
        plt.ylabel("Least Squares/PCA Slope of Best Fit Line")  
        plt.legend()  
        plt.savefig(f"figures/{title}.png", format="png")  
        plt.close()
```

```
[278]: problem1cd("problem1c", add_x_noise=False)
```

```
[279]: problem1cd("problem1d", add_x_noise=True)
```

4 Problem 2

```
[352]: class Globals:  
        DATA_FOLDER = "data"  
  
        def get_genomic_matrix_data(use_mode=True):  
            """Returns the binary X matrix as defined in the problem statement."""  
            datapath = os.path.join(Globals.DATA_FOLDER, "p4dataset2020.txt")  
            with open(datapath) as f:  
                lines = f.readlines()
```

```

assert len(lines) == 995
cleaned = [[x.strip() for x in line.split(" ")]
            for line in lines]
pop = np.array([line[2] for line in cleaned])
gender = np.array([int(line[1]) for line in cleaned])
ids = np.array([line[0] for line in cleaned])

if use_mode:
    mapping = {'A': 0, 'C': 1, 'T': 2, 'G': 4}
    data = np.array([[mapping[x] for x in line[3:]]
                     for line in cleaned])

    # Find mode.
    mode = scipy.stats.mode(data, axis=0).mode
    X = (data != mode).astype(int)
else:
    mapping = {'A': [0, 0, 0, 1], 'C': [0, 0, 1, 0], 'T': [0, 1, 0, 0], 'G':
→ [1, 0, 0, 0]}
    data = np.array([[mapping[x] for x in line[3:]]
                     for line in cleaned])
    X = data.reshape((data.shape[0], -1))
return X, pop, gender, ids

```

```

[325]: def problem2b(use_mode):
    X, pop, _, _ = get_genomic_matrix_data(use_mode)
    X = X - np.mean(X, axis=0)
    pca_model = decomposition.PCA(n_components=2)
    pca_model = pca_model.fit(X)
    v1, v2 = pca_model.components_
    v1, v2 = v1 / np.linalg.norm(v1), v2 / np.linalg.norm(v2)
    xs, ys = np.dot(X, v1).flatten(), np.dot(X, v2).flatten()
    # xs, ys = projections[:,0], projections[:,1]
    label_map = {
        'ACB': 'African Caribbeans in Barbados',
        'ASW': 'Americans of African Ancestry in SW USA',
        'ESN': 'Esan in Nigeria',
        'GWD': 'Gambian in Western Divisions in the Gambia',
        'LWK': 'Luhya in Webuye, Kenya',
        'MSL': 'Mende in Sierra Leone',
        'YRI': 'Yoruba in Ibadan, Nigeria',
    }
    for label in np.unique(pop):
        idx = (pop == label).nonzero()
        plt.scatter(xs[idx], ys[idx], label=label_map[label])
    plt.legend()
    plt.title("Genomic sequence data project into 2 principal components")
    plt.xlabel("v1")
    plt.ylabel("v2")

```

```
plt.savefig(f"figures/genomic_2d_projection_use_mode={use_mode}.png",
↪format="png")
plt.close()
```

```
[394]: problem2b(use_mode=True)
```

```
[326]: def problem2d():
    X, pop, gender, _ = get_genomic_matrix_data()
    X = X - np.mean(X, axis=0)
    pca_model = decomposition.PCA(n_components=3)
    pca_model = pca_model.fit(X)
    v1, _, v3 = pca_model.components_
    v1, v3 = v1 / np.linalg.norm(v1), v3 / np.linalg.norm(v3)
    xs, ys = np.dot(X, v1).flatten(), np.dot(X, v3).flatten()
    # xs, ys = projections[:,0], projections[:,1]
    label_map = {
        1: 'Male',
        2: 'Female'
    }
    for label in np.unique(gender):
        idx = (gender == label).nonzero()
        plt.scatter(xs[idx], ys[idx], label=label_map[label])
    plt.legend()
    plt.title("Genomic sequence data project into 1st and 3d principal
↪components")
    plt.xlabel("v1")
    plt.ylabel("v3")
    plt.savefig("figures/genomic_2d_projection_2.png", format="png")
    plt.close()
```

```
[252]: problem2d()
```

```
[327]: def problem2f():
    X, pop, gender, _ = get_genomic_matrix_data()
    X = X - np.mean(X, axis=0)
    pca_model = decomposition.PCA(n_components=3)
    pca_model = pca_model.fit(X)
    _, _, v3 = pca_model.components_
    v3 = v3.flatten() / np.linalg.norm(v3)

    plt.title("Magnitude of 3rd Principal Component")
    plt.xlabel('Nucleobase Index')
    plt.ylabel('Absolute Value')
    plt.plot(range(len(v3)), np.abs(v3))
    plt.savefig("figures/x_y_chromosomes.png", format="png")
    plt.close()
```

```
[270]: problem2f()
```

```
[314]: def problem2h():  
        problem2b(use_mode=False)
```

```
[315]: problem2h()
```

```
[395]: def read_metadata():  
        """Returns a dictionary of metadata, M.  
  
        M[property][I_ID] returns the corresponding property of I_ID.  
        """  
        metadata = pd.read_csv(os.path.join(Globals.DATA_FOLDER, "20130606_g1k.  
↪ped"),  
                                header=0, sep='\t').set_index('Individual ID')  
        metadata = metadata.to_dict()  
        return metadata  
  
def problem2i(color_label):  
    meta = read_metadata()  
    X, _, _, ids = get_genomic_matrix_data(use_mode=False)  
    X = X - np.mean(X, axis=0)  
    pca_model = decomposition.PCA(n_components=4)  
    pca_model = pca_model.fit(X)  
    _, _, v3, v4 = pca_model.components_  
    xs, ys = np.dot(X, v3).flatten(), np.dot(X, v4).flatten()  
    # xs, ys = projections[:,0], projections[:,1]  
  
    labels = meta[color_label]  
    data_labels = np.array([labels[_id] for _id in ids])  
    for label in np.unique(data_labels):  
        idx = (data_labels == label).nonzero()  
        plt.scatter(xs[idx], ys[idx], label=label)  
    plt.legend()  
    plt.title("Genomic sequence data project into 1st and 3d principal_  
↪components")  
    plt.xlabel("v1")  
    plt.ylabel("v4")  
    plt.savefig("figures/genomic_2i_projection_4th.png", format="png")  
    plt.close()
```

```
[396]: problem2i(color_label='Relationship')
```